

Arya Chakraborty [22MSD7020]

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
```

```
seeds = pd.read_excel("Pumpkin_Seeds_Dataset.xlsx")
```

```
seeds.head()
```

	Area	Perimeter	Major_Axis_Length	Minor_Axis_Length	Convex_Area
0	56276	888.242	326.1485	220.2388	56831
1	76631	1068.146	417.1932	234.2289	77280
2	71623	1082.987	435.8328	211.0457	72663
3	66458	992.051	381.5638	222.5322	67118
4	66107	998.146	383.8883	220.4545	67117


	Equiv_Diameter	Eccentricity	Solidity	Extent	Roundness
0	267.6805	0.7376	0.9902	0.7453	0.8963
1.4809					
1	312.3614	0.8275	0.9916	0.7151	0.8440
1.7811					
2	301.9822	0.8749	0.9857	0.7400	0.7674
2.0651					
3	290.8899	0.8123	0.9902	0.7396	0.8486
1.7146					
4	290.1207	0.8187	0.9850	0.6752	0.8338
1.7413					

	Compactness	Class
0	0.8207	Çerçvelik
1	0.7487	Çerçvelik
2	0.6929	Çerçvelik
3	0.7624	Çerçvelik
4	0.7557	Çerçvelik

```
sns.heatmap(seeds.corr(), annot=True, cmap="YlGnBu")
plt.title("Correlation Heatmap", fontsize=16)
plt.show()
```


C:\Users\chakr\AppData\Local\Temp\ipykernel_13280\3006396044.py:1:
FutureWarning: The default value of numeric_only in DataFrame.corr is
deprecated. In a future version, it will default to False. Select only
valid columns or specify the value of numeric_only to silence this
warning.

```
sns.heatmap(seeds.corr(), annot=True, cmap="YlGnBu")
```



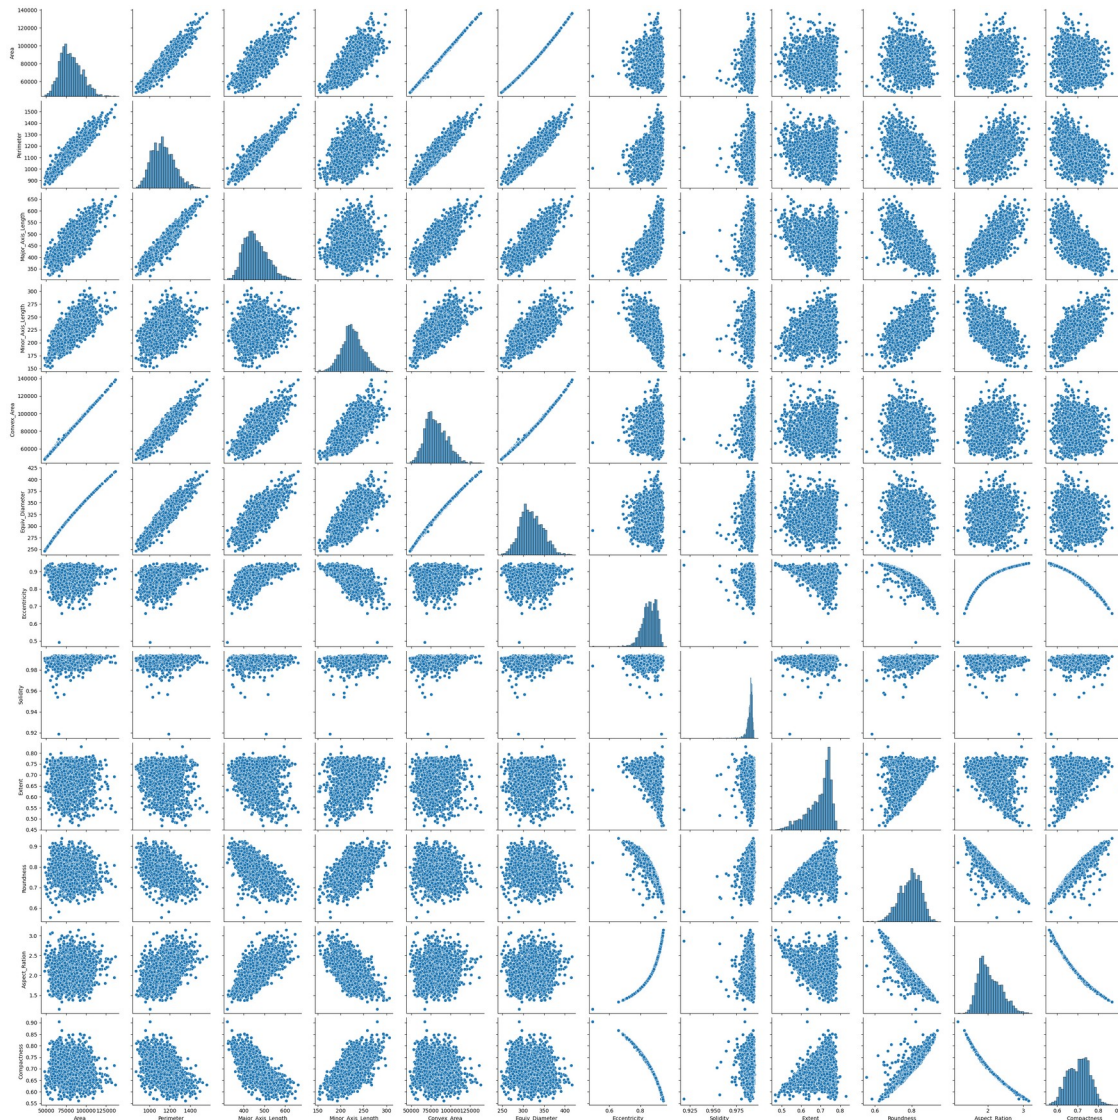
```
sns.pairplot(seeds, kind='box')
```

<seaborn.axisgrid.PairGrid at 0x2c2edbd2c50>



```
sns.pairplot(seeds)
```

```
<seaborn.axisgrid.PairGrid at 0x2c2ede241f0>
```



```
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
```

```
X = seeds[["Area", "Perimeter",
           "Major_Axis_Length", "Minor_Axis_Length", "Convex_Area",
           "Equiv_Diameter", "Eccentricity", "Solidity", "Extent", "Roundness",
           "Aspect_Ration", "Compactness"]]
```

```
y=seeds['Class']
```

```
scaler = StandardScaler()
X_scaled=scaler.fit_transform(X)
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled,y,
test_size=0.3, random_state=0)
```

```
svm = SVC(kernel='linear')
svm.fit(X_train, y_train)
```

```

SVC(kernel='linear')

# Predict the labels of the test set
y_pred = svm.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

Accuracy: 0.8533333333333334

from sklearn.metrics import classification_report, confusion_matrix

confusion_matrix(y_test, y_pred)

array([[344, 45],
       [ 65, 296]], dtype=int64)

print(classification_report(y_test, y_pred))

```

	precision	recall	f1-score	support
Çerçvelik	0.84	0.88	0.86	389
Ürgüp Sivrisi	0.87	0.82	0.84	361
accuracy			0.85	750
macro avg	0.85	0.85	0.85	750
weighted avg	0.85	0.85	0.85	750

HyperParameter Tuning

```

from sklearn.model_selection import GridSearchCV
param_grid = {'C': [0.1, 1, 10], 'kernel': ['linear', 'rbf'], 'gamma':
['scale', 'auto']}

svm1 = SVC()
grid_search = GridSearchCV(svm1, param_grid, cv=5, scoring='accuracy',
n_jobs=-1)

grid_search.fit(X_train, y_train)

GridSearchCV(cv=5, estimator=SVC(), n_jobs=-1,
param_grid={'C': [0.1, 1, 10], 'gamma': ['scale',
'auto'],
'kernel': ['linear', 'rbf']},
scoring='accuracy')

print("Best parameters:", grid_search.best_params_)
print("Best score:", grid_search.best_score_)

```

```
Best parameters: {'C': 1, 'gamma': 'scale', 'kernel': 'rbf'}  
Best score: 0.8942857142857144
```

```
best_model = grid_search.best_estimator_  
y_pred = best_model.predict(X_test)  
accuracy = accuracy_score(y_test, y_pred)  
print("Test accuracy:", accuracy)
```

```
Test accuracy: 0.8626666666666667
```

Conclusion : we can see that during cross validation the score obtained is .89 and the final test score is .86, as the difference is not that big, we can consider that there are no significant overfitting problem present in the model.

To optimize further we can use Bayesian optimization or Random Search CV to see if we can get better result. But considering these best fitted models we can clearly conclude that we have reached the threshold for the parameter tuning.